

# Computer Vision HW2 Report

## 1. Implementation

I implemented a RANSAC class, these are the parameters, the usage of each would be described in later section.

```
21 class RANSAC:
22     def __init__(self, threshold=0.75, iter_num=10000, ransac_tolerance=3, image_list=None, early_termination_ratio=0.9):
23         self.threshold = threshold
24         self.iter_num = iter_num
25         self.ransac_tolerance = ransac_tolerance
26         self.h = None
27         self.image_list = image_list
28         self.sift = cv2.SIFT_create()
29         self.early_termination_ratio = early_termination_ratio
```

At the start of the process, it would set the middle image as the base (to get a better result), and stitch other images from middle to left and from middle to right.

```
178     def __call__(self):
179         mid = len(self.image_list) // 2
180         base = self.image_list[mid]
181         for image in self.image_list[mid - 1::-1]:
182             base = self.sift_and_stitch(base, image, base_side: "right")
183
184         for image in self.image_list[mid + 1:]:
185             base = self.sift_and_stitch(base, image, base_side: "left")
186
187         return base
```

In the *sift\_and\_stitch* function, it firstly calculates the proper neighbors by *find\_knn*, then use *ransac* to sample the best homography matrix.

```
154     def sift_and_stitch(self, base_img, addition_img, base_side="right"):
155         h1, w1 = base_img.shape[:2]
156         h2, w2 = addition_img.shape[:2]
157         pts_base, pts_addition = self.find_knn(base_img, addition_img)
158
159         self.ransac(pts_addition, pts_base)
```

After calculating the homography matrix, it transforms the corners of addition image by the homography and get the transformed positions. Using the transformed corners and the corners of the base image, we can calculate the needed image

size  $(max\_x - min\_x, max\_y - min\_y)$  and the affine translation matrix  $\begin{bmatrix} 1 & 0 & -min\_x \\ 0 & 1 & -min\_y \\ 0 & 0 & 1 \end{bmatrix}$ .

```

161 base_corners = np.array( object: [[0, 0], [0, h1], [w1, h1], [w1, 0]], dtype=float)
162 addition_corners = np.array( object: [[0, 0], [0, h2], [w2, h2], [w2, 0]], dtype=float)
163 addition_corners = cv2.perspectiveTransform(addition_corners.reshape(1, -1, 2), self.h).reshape(-1, 2)
164
165 min_x, min_y = np.min(np.vstack([base_corners, addition_corners]), axis=0)
166 max_x, max_y = np.max(np.vstack([base_corners, addition_corners]), axis=0)
167 new_size = (int(max_x - min_x), int(max_y - min_y))
168
169 affine = np.array( object: [[1, 0, -min_x], [0, 1, -min_y], [0, 0, 1]], dtype=float)

```

Then we can transform the images by *warpPerspective* and stitch them.

```

171 img_after = cv2.warpPerspective(addition_img, affine @ self.h, new_size)
172 base_img_after = cv2.warpPerspective(base_img, affine, new_size)
173
174 stitched = self.stitch(base_img_after, img_after, base_side=base_side)
175
176 return stitched

```

In the *find\_knn* function, it calculates the keypoints and for each description in image 1, calculate all the distance of it and the descriptors in image 2, get the keypoints with the two lowest distances. If the distance of the first closest is smaller than *threshold* \* the distance of the second closest, then record the positions to be possible matches.

```

31 def find_knn(self, img1, img2):
32     kp1, des1 = self.sift.detectAndCompute(cv2.cvtColor(img1, cv2.COLOR_BGR2GRAY), None)
33     kp2, des2 = self.sift.detectAndCompute(cv2.cvtColor(img2, cv2.COLOR_BGR2GRAY), None)
34     pts1, pts2 = [], []
35
36     for kp, des in zip(kp1, des1):
37         matches = [(kp_, np.linalg.norm(des - des_)) for kp_, des_ in zip(kp2, des2)]
38         matches.sort(key=lambda x: x[1])
39         if matches[0][1] < self.threshold * matches[1][1]:
40             pts1.append(kp.pt)
41             pts2.append(matches[0][0].pt)
42
43     return pts1, pts2

```

In the *ransac* function, it runs for *iter\_num* iterations, and sample four random indices as the position pairs to fit homography, and calculate the homography by the *calc\_h* function. Then transform all the points on addition image to its position on base image.

```

45     def ransac(self, pts1, pts2):
46         max_in = 0
47         best_h = None
48
49         n = len(pts1)
50         for _ in range(self.iter_num):
51             idx = random.sample(range(n), k=4)
52             pts1_ = np.array([pts1[i] for i in idx])
53             pts2_ = np.array([pts2[i] for i in idx])
54             h = self.calc_h(pts1_, pts2_)
55
56             pts1__ = np.append(pts1, np.ones((n, 1)), axis=1)
57
58             pts_h = h @ pts1__.T
59             pts_h = pts_h / pts_h[2]
60             pts_h = pts_h[:2].T

```

After transforming the points, compare the transformed position and its corresponding point on base image, calculate the numbers of the close enough point pairs. If this sample is better, then record it. If the homography is great enough, then break the process.

```

62         in_count = np.sum(np.linalg.norm(pts_h - pts2, axis=1) < self.ransac_tolerance)
63
64         if in_count > max_in:
65             max_in = in_count
66             best_h = h
67
68         if max_in / n > self.early_termination_ratio:
69             break
70
71         self.h = best_h

```

To calculate homography, I fill the A matrix and b vector, then solve the homography parameters by least square.

$$\begin{array}{c} \mathbf{A} \end{array}
 \begin{bmatrix} x_1 & y_1 & 1 & 0 & 0 & 0 & -x_1\hat{x}_1 & -y_1\hat{x}_1 \\ x_2 & y_2 & 1 & 0 & 0 & 0 & -x_2\hat{x}_2 & -y_2\hat{x}_2 \\ x_3 & y_3 & 1 & 0 & 0 & 0 & -x_3\hat{x}_3 & -y_3\hat{x}_3 \\ x_4 & y_4 & 1 & 0 & 0 & 0 & -x_4\hat{x}_4 & -y_4\hat{x}_4 \\ 0 & 0 & 0 & x_1 & y_1 & 1 & -x_1\hat{y}_1 & -y_1\hat{y}_1 \\ 0 & 0 & 0 & x_2 & y_2 & 1 & -x_2\hat{y}_2 & -y_2\hat{y}_2 \\ 0 & 0 & 0 & x_3 & y_3 & 1 & -x_3\hat{y}_3 & -y_3\hat{y}_3 \\ 0 & 0 & 0 & x_4 & y_4 & 1 & -x_4\hat{y}_4 & -y_4\hat{y}_4 \end{bmatrix}
 \begin{array}{c} \mathbf{h} \end{array}
 \begin{bmatrix} h_{11} \\ h_{12} \\ h_{13} \\ h_{21} \\ h_{22} \\ h_{23} \\ h_{31} \\ h_{32} \end{bmatrix}
 = h_{33}
 \begin{array}{c} \mathbf{b} \end{array}
 \begin{bmatrix} \hat{x}_1 \\ \hat{x}_2 \\ \hat{x}_3 \\ \hat{x}_4 \\ \hat{y}_1 \\ \hat{y}_2 \\ \hat{y}_3 \\ \hat{y}_4 \end{bmatrix}$$

```

73     @staticmethod
74     def calc_h(pts1, pts2):
75         a = np.zeros((8, 8))
76         b = np.zeros(8)
77
78         a[:4, :2] = pts1[:4]
79         a[:4, 2] = 1
80         a[4:, 3:5] = pts1[:4]
81         a[4:, 5] = 1
82         a[:4, 6] = -pts1[:4, 0] * pts2[:4, 0]
83         a[:4, 7] = -pts1[:4, 1] * pts2[:4, 0]
84         a[4:, 6] = -pts1[:4, 0] * pts2[:4, 1]
85         a[4:, 7] = -pts1[:4, 1] * pts2[:4, 1]
86
87         b[:4] = pts2[:4, 0]
88         b[4:] = pts2[:4, 1]
89
90         h = np.linalg.lstsq(a, b, rcond=None)[0]
91         h = np.append(h, values: 1).reshape(3, 3)
92         return h

```

In the *stitch* function, it firstly converts the images to gray and calculate masks from the gray images. For the overlapped origin, calculate another overlap mask to be blended in following process.

```

139     def stitch(self, base_img, addition_img, base_side="right"):
140         gray_base = cv2.cvtColor(base_img, cv2.COLOR_BGR2GRAY)
141         gray_addition = cv2.cvtColor(addition_img, cv2.COLOR_BGR2GRAY)
142
143         mask_base = self.calc_mask(gray_base)
144         mask_addition = self.calc_mask(gray_addition)
145         mask_overlap = mask_base & mask_addition
146
147         if base_side == "right":
148             img_result = self.no_blending(addition_img, base_img, mask_addition, mask_base, mask_overlap)
149         else:
150             img_result = self.no_blending(base_img, addition_img, mask_base, mask_addition, mask_overlap)
151
152         return img_result

```

For blending, I've implemented three functions to blend the result. First, *no\_blending* simply paste the image to the left on the image to the right.

```

102     @staticmethod
103     def no_blending(img_left, img_right, mask_left, mask_right, mask):
104         img_result = img_right.copy()
105         img_result[mask_left] = img_left[mask_left]
106         img_result[mask_left & mask_right] = img_left[mask_left & mask_right]
107
108         return img_result

```

Second, *weighted\_blending* use *cv2.addWeighted* to average the two images and replace the overlapped part with the blended result.

```

131     @staticmethod
132     def weighted_blending(img_left, img_right, mask_left, mask_right, mask):
133         img_result = img_right.copy()
134         img_result[mask_left] = img_left[mask_left]
135         img_result[mask_left & mask_right] = cv2.addWeighted(img_left, alpha=0.5, img_right, beta=0.5, gamma=0)[mask_left & mask_right]
136
137         return img_result

```

Third, *linear\_blending* take the leftmost and rightmost pixel of overlapped region, and from the left to the right, the result takes weighted from the left image and (1-weighted) from the right image. For the weights, it is gradually decreased from the leftmost overlapped pixel to the rightmost overlapped pixel.

```

110     @staticmethod
111     def linear_blending(img_left, img_right, mask_left, mask_right, mask):
112         leftmost = np.min(np.where(mask)[1])
113         rightmost = np.max(np.where(mask)[1])
114
115         step = 1 / (rightmost - leftmost)
116         alpha_mask = np.zeros_like(mask, dtype=float)
117
118         img_result = img_right.copy()
119         img_result[mask_left] = img_left[mask_left]
120
121         for i in range(leftmost, rightmost):
122             alpha_mask[:, i] = (i - leftmost) * step
123
124         alpha_mask_3d = alpha_mask[..., np.newaxis]
125         one_minus_alpha_mask_3d = 1 - alpha_mask_3d
126
127         img_result[mask] = img_left[mask] * one_minus_alpha_mask_3d[mask] + img_right[mask] * alpha_mask_3d[mask]
128
129         return img_result

```

For *calc\_mask*, I noticed that border of transformed image could be some dark pixel but can't simply capture that by *gray\_image*  $\neq 0$ , so I implement another function to extract the pixels with its neighbors all greater than zero as the mask.

(the difference is showed below)

```

94     @staticmethod
95     def calc_mask(img):
96         mask = img  $\neq$  0
97         direction = np.array([[0, 1], [1, 0], [0, -1], [-1, 0]])
98         for d in direction:
99             mask  $\&=$  np.roll(img, d, axis=(0, 1))  $\neq$  0
100         return mask

```

## 2. Stitching result (linear blending)



### 3. Comparison

1. No blending (custom defined mask):



We can find some discontinuous part in the result and also the color is not blended.






2. Weighted blending (custom defined mask):



The color is blended better, but still have some obvious color gap.



3. Linear blending (custom defined mask):





I think the only flaw is the bottom part.

Extra: Linear blending ( $gray\_img \neq 0$  mask):



It is obvious that there are some black lines because of the dark border of the transformed images.

